

Modal logics, temporal and dynamic

Adam Gashlin

Contents

- [Propositional logic](#)
- [Modal logic](#)
- [Temporal logic](#)
- [Dynamic logic](#)

Overview

This report deals with a branch of formal logic called Modal Logic, in which we are concerned with the truth of propositions in different situations, and how those situations relate to each other logically. Example situations include points in time, conceivable worlds, believable knowledge, and states of a computation.

There is a good deal of theoretical mechanism that applies equally well to many kinds of Modal Logic, such as the unifying Kripke semantics. The essential core of Kripke semantics is a binary relation on the situational "worlds", which defines how propositions in one world may be referenced by propositions in another. This report examines in detail two particular families of modal logic, Temporal Logic and Dynamic Logic.

In temporal logic truth is given for particular points in time. Times are related by " s is in the past of t " and " s is in the future of t ". We can express statements with reference to some or all times, future or past, such as " p or q will be true" or "if r was always true, r is true". These are useful, for instance, for verifying the eventual proper behavior of programs, invariants such as "if the operating system gets a request, the request will be granted".

In dynamic logic, truths hold for states of an executing program, and the relations are the input/output relation of each program. This lets us express such statements as "when program a finishes, p is true" and "if p or q are true, then when program b finishes, q is true". This can be useful for formally specifying correctness properties, such as "if a is positive before the loop, a is zero after the loop terminates".

Propositional logic

Formal logic

Our concern in a formal logic is to abstract away every detail that is not necessary for preserving the validity of statements. Much reasoning is valid on a formal level: a properly formed argument is assured to be valid by virtue of its form. This abstraction allows logicians to study the structure of logic apart from dealing with the specifics of arguments, and it encourages attention to the essential details of a situation.

Most of the operations of formal logic can be executed mechanically, to varying degrees of effectiveness. There can be great practical value in expressing a problem in the terms of a logic, if there are efficient automatic methods available.

Propositions

A *proposition* is simply a statement which we regard as true or false. An *atomic proposition* is defined to have no internal structure that we are interested in.

Consider the two atomic propositions

- p = "My pants are blue"
- q = "My hair is black"

By calling these atomic, we are ignoring as irrelevant any common details, such as that both of these statements are about me and involve colors. An atomic proposition only has an identity, like the letters p and q here, which lets us know whether propositions can be considered identical.

Connectives

In propositional logic, we reduce statements to atomic propositions, joined by a small family of connectives into larger propositions. If we wanted a proposition that means "my pants are blue and my hair is black", we might introduce a new atomic proposition called s , but this would be a poor choice. For instance, it would allow us to make the statement " s , but not p ". This is never true given the earlier definition of p (it is a *contradiction*), but that is not clear from the form of the statement.

Instead of s , we can call this proposition $p \wedge q$. The propositional connective \wedge in the middle is read as "and", and the combined proposition is the statement "both p and q are true".

By using this standard method of connecting existing propositions to name a new one, we get a name that is much more useful. If we know that both proposition p and proposition q are true, we automatically know that proposition $p \wedge q$ is true, without having to look up any details about what these propositions mean. Likewise the contradiction from earlier would be expressed as " $p \wedge q$ but not p ", in which the contradiction is immediately apparent.

Implication

An important connective is "implies", written as \rightarrow . This is used to write propositions such as "if it is noon, the sun is up". Note that implication doesn't say that two statements are always the same; it is quite possible for the sun to be up when it isn't noon. The only strong claim made by $p \rightarrow q$ is that *when* p is true, q is also true. Therefore $p \rightarrow q$ is only false when p is true and q is false.

Parenthesis

Consider a proposition $a \vee b \wedge c$. How is this to be interpreted? We could consider the whole proposition to be of the form $a \vee X$, where the X is $b \wedge c$. On the other hand it seems that we could just as well consider it to be $X \wedge c$ where X is $a \vee b$. These interpretations give different meanings to the proposition, for instance when a = true, b = true, c = false, the first interpretation is true and the second is false.

We eliminate these issues of ambiguity by adding parenthesis around parts of the expression that are to be considered as a unit. For instance, if we had instead written $(a \vee b) \wedge c$, we would be insisting on an interpretation where the proposition is of the form $X \wedge c$, where X is $(a \vee b)$. The alternative $a \vee X$ wouldn't make sense: for one thing $a \vee X$ doesn't have a place for the left parenthesis, for another the X would have to

be $b) \wedge c$, with a lonely right parenthesis.

To summarize, we put a pair of parenthesis around a proposition to "wrap" it, so that any interpretation has to "unwrap" it as a whole rather than breaking it up and giving parts to surrounding expressions.

As a side note, parenthesis are only necessary because we are using *infix* notation, where connectives are written in between the propositions they connect, like $a \vee b$. An alternative is *postfix* notation, where the connectives are written afterwards, such as $a b \vee$. This form is unambiguous, as the connectives always connect the two previous units in the expression. For example the two different interpretations of $a \vee b \wedge c$ are written in postfix as $a b c \vee \wedge$ (for $a \vee (b \wedge c)$) and $a b \vee c \wedge$ (for $(a \vee b) \wedge c$). The problem is that postfix is considered unintuitive and difficult to read, so we stick with infix.

Evaluation

We haven't said yet what color my hair or pants actually are. There are some statements whose truth doesn't depend on the truth of the atomic propositions, for instance contradictions are always false and their opposite, *tautologies*, are always true. But for everything else, the truth value of a proposition depends on the truth values of the atomic propositions it is made from.

The process of determining the truth value is *evaluation*, and it proceeds according to the structure of a statement. For a statement like the $p \wedge q$ we have been considering, consult the *truth table* for \wedge below.


\wedge			\vee			\rightarrow			\neg	
φ	ψ	$\varphi \wedge \psi$	φ	ψ	$\varphi \vee \psi$	φ	ψ	$\varphi \rightarrow \psi$	φ	$\neg \varphi$
T	T	T	T	T	T	T	T	T	T	F
T	F	F	T	F	T	T	F	F	F	T
F	T	F	F	T	T	F	T	T		
F	F	F	F	F	F	F	F	T		

The Greek letters φ (phi) and ψ (psi) here are used to indicate any proposition, not just an atomic proposition. In $p \wedge q$, $\varphi = p$ and $\psi = q$. If φ is **True** and ψ is **False**, $\varphi \wedge \psi$ is **False**.


The benefit of this method over saying " \wedge is 'and'" is that the table leaves no room for interpretation. While there may not be many ways of interpreting "and", consider "or", which is written as \vee in propositional logic, and which has the truth table shown above.

In English it is not unusual for "x or y" to exclude the case where both x and y are true, but the truth table tells us unambiguously that φ **True** and ψ **True** makes $\varphi \vee \psi$ **True**.

Consider a more complicated statement such as $(\neg p \wedge q) \vee (p \wedge \neg q)$. First off, $\neg q$ means "not q", which has the truth table above. At this point it is helpful to think of each of the atomic propositions p and q as simply something that can be true or false. If we say that p is true and q is false, this forms our *model* for the state of the world, our model is formed from the *valuation* " p is true, q is false". To evaluate $(\neg p \wedge q) \vee (p \wedge \neg q)$, we break it down until we arrive at statements whose truth can be determined directly from the model, then we build back up to the full statement. This can be pictured as building a tree



and filling it with truth values.



Thus for this model, with the valuation "p is true, q is false", $(\neg p \wedge q) \vee (p \wedge \neg q)$ is true.

Evaluation by definition of $M \models \varphi$

Formally, we treat evaluation as extending the valuation to cover all possible propositions formed from the atomic propositions. If a proposition φ is true in model M , we write $M \models \varphi$.

$M \models p$	if p is a true atomic proposition
$M \models \neg\varphi$	if $M \models \varphi$ is not true
$M \models \varphi \wedge \psi$	if $M \models \varphi$ and $M \models \psi$
$M \models \varphi \vee \psi$	if one or both of $M \models \varphi$ and $M \models \psi$

Note that only one of these rules applies for every possible proposition, so if the left hand side matches the form of a proposition but the condition on the right hand side doesn't hold, then the proposition is false. If $M \models \varphi$ is false, we write $M \not\models \varphi$.

Modal logic

Multiple states

So far everything discussed holds for classical propositional logic as well as modal logic. The first place where modal logic differs is that a model can have multiple *states*, a.k.a. "situations" or "worlds", whereas classical propositional logic considers only one set of truths at a time.

As an example, say we have one state, call it s , where the atomic propositions p and q are both true. We also have a state t , where only p is true and q is false. In modal logic a model M has at least two components: the set of states S_M and a valuation V_M . In the example, $S_M = \{s, t\}$, the two states, and V_M is a function that gives the atomic propositions that are true in a specified state: $V_M(s) = \{p, q\}$ and $V_M(t) = \{p\}$.

To show that we are discussing truths in a particular state s , we write $M \models_s$ instead of simply $M \models$. In the running example, $M \models_s p$ and $M \models_s q$, as well as $M \models_t p$. However it is not the case that $M \models_t q$, so we write $M \not\models_t q$. We can also use the earlier definition of \models to extend \models_s and \models_t to larger propositions, for instance $M \models_s p \wedge q$ and $M \models_t p \wedge \neg q$.

Often we want to reason without knowing what state is the "real world". If we find a proposition that is true across all states, we can use it without reference to a particular state. If this is so we say the proposition is *true in the model*, and we write $M \models \varphi$ without a subscript. In the example we have $M \models p$ and $M \models q \vee p$, for instance.

Referring to other states

The real power of modal logic comes from being able to make reference to other states. For this purpose we add two symbols to our proposition language: \Box ("box") and \Diamond ("diamond"). These are used to prefix other propositions to modify their interpretation. $\Box\varphi$ means that in all related states (to be defined in the next section), φ is true, whereas $\Diamond\varphi$ asserts that there is at least one related state where φ is true.

Note that $\Box\varphi$ and $\Diamond\varphi$ are only about related states, they make no reference to what is true in the current state. To claim that φ is true in the current state and in all related states, use the proposition $\varphi \wedge \Box\varphi$. Also note that $\Box\varphi$ can be true even if there are no states where φ is true; this is the case when there are no related states.

We can formally define \Box and \Diamond by adding the following rules to our definition of \models :

$M \models_s \Box \varphi$ if for all states t where $sR_M t$, $M \models_t \varphi$
 $M \models_s \Diamond \varphi$ if there is a state t such that $sR_M t$ and $M \models_t \varphi$

$sR_M t$ means that s is related to t by relation R_M , which we will explain in the next section.

Related states

The meaning of \Box and \Diamond in a given model is specified by a binary relation between pairs of states. A *binary relation* is simply a set of ordered pairs of objects. The term "binary" here means that the relation deals with two objects at a time, i.e. it is about pairs. If a relation R contains a pair of states, then those states are considered to be related by R . For example the relation $R = \{(s, t), (s, u), (t, u), (u, t)\}$ means that s is related to t by R , which we write as sRt . The converse is not true, however: it is not the case that tRs (as there is no pair (t, s)). Both tRu and uRt are true.


One way to understand the meaning of a relation in a model is that it describes the set of states that are accessible, by \Box and \Diamond , from a given state. $M \models_s \Box q$ only refers to the truth of q in those states t that are related to s by $sR_M t$.

More on this in the examples below.

An example

For a simple example (based on Harel et al. exercise 5.1), consider a model with the states S_M , relation R_M , and truth values determined by V_M :

$$\begin{aligned}
 S_M &= \{s, t, u\}, \\
 R_M &= \{(s, t), (s, u), (t, u), (u, t)\}, \\
 V_M(s) &= \{\}, V_M(t) = \{p, q\}, \\
 V_M(u) &= \{p\}.
 \end{aligned}$$



This is depicted in the diagram on the right, where circled letters are states, arrows show relations, and colored shapes enclose states where the like-colored propositions are true.

Here, we have $M \models_x \Box p$ for any state x , as all states y that x is related to (i.e. that it has an arrow from x to y) have $M \models_y p$. Therefore we can write $M \models \Box p$.

We also have $M \models_s \Diamond q$ and $M \models_u \Diamond q$, but not $M \models_t \Diamond q$.

$M \models_t \Diamond \Diamond q$ is true, as there is some state (u) related from t that relates to some state (t again) where q is true.

Frames

A proposition can be true in one state of a model, or it can be true across all states in the model. Then there are tautologies such as $p \vee \neg p$, which are true regardless of model, as they are true regardless of the truth values of the atomic propositions. In between propositions that are and that are true in a model, there is another level of propositions that are *valid in a frame*.

1. φ is true in a state s in a model M , $M \models_s \varphi$
2. φ is true in a model M , true in all states of the model, $M \models \varphi$

3. φ is valid in a frame F , true on all models on a frame, $F \models \varphi$
4. φ is a tautology, valid regardless of frame, $\models \varphi$

A *frame* consists of just the set of states S_M and the relation R_M used to construct the model, that is it does not specify what the truth values actually are in the states. The propositions that are valid on a frame thus depend on the properties of the relation. Some examples follow (from Goldblatt, p. 12, theorem 1.12).

Reflexivity

If a relation says that a state is always related to itself (sRs for all states s , the relation is called *reflexive*. In any frame that has a reflexive relation, all propositions of the form $\Box\varphi \rightarrow \varphi$ will be valid. If all related states have φ true, then this state must have φ true as well, since it is related to itself.


Symmetry

In a symmetric relation R , if sRt then tRs , so all relations go both directions. In a frame with a symmetric relation, all propositions of the form $\varphi \rightarrow \Box\Diamond\varphi$ are true. That is, if φ is true at s , then in every state t related to s , there is some related state u where φ is true. To see why this is true, u can just be s since both sRt and tRs .

Transitivity

A transitive relation is always "inherited" through a middle state. If sRt and tRu , then, if R is a transitive relation, it is also the case that sRu . This allows the middle state to be cut out so that the relation applies directly between two states that are otherwise known to be second-degree related.

In a frame with a transitive relation, all propositions with the form $\Box\varphi \rightarrow \Box\Box\varphi$ are valid. This is because, as shown on the right, all states that are second-degree related (such as u) are also first-degree related (because of transitivity), thus if all first-degree related states have φ (i.e. $\Box\varphi$), then all second-degree related states have φ (i.e. $\Box\Box\varphi$).



Seriality

A serial relation is usually on an infinite set, like a set of points in time that we assume will go on for ever; seriality is the property that any state always has a related state. In a frames with a serial relation, all propositions with the form $\Box\varphi \rightarrow \Diamond\varphi$ are valid. This is simply the statement that we don't have any dead ends where there are no related states, where $\Box\varphi$ can be true without φ actually being true anywhere (because "all related states" is nothing).

Temporal logic

A time-dependent logic

Temporal logic is a branch of modal logic which supports propositions that are true at one time yet are false at another. Moreover, because this is a full-featured modal logic, the propositions true at some time can make reference to other times. Thus we can give the proposition p for "the machine is working", $\Diamond p$ for "at some point in the future the machine will work", and $\Box p$ for "the machine will work at all times in the future".

In a temporal logic model M , we consider the states of the model S_M to be points in time, and the relation R_M is defined between two points in time, one of which is before the other. $sR_M t$ means s is before t , therefore t

is in the future of s . Those propositions which are true at a time s are given by $V_M(s)$, as in modal logic in general.

Properties of the temporal relation

Let's consider a few of the [properties of relations](#) mentioned above, and whether they should apply to our model of time. It is normally the case that we do not want to consider the present as part of the future, which is to say the relation R_M is not reflexive, in fact it is *irreflexive* as no time relates to itself. We also insist that times only have one order, so $sR_M t$ and $tR_M s$ can't both be true. This makes R_M not symmetric, indeed it is a combination of the irreflexive property and the *antisymmetric* property, which states that $sR t$ and $tR s$ only if $s = t$. These properties, irreflexivity and antisymmetry, do not correspond to valid propositions across frames, unlike reflexivity and symmetry do (Goldblatt p. 14).

We probably also want R_M to have seriality, to ensure that time goes on forever.

Transitivity is an important property that R_M must have if M is to make sense as a model of time. A transitive relation means that $sR_M t \wedge tR_M u \rightarrow sR_M u$, so a point in the future of the future is in the future of the present. It would make little sense that a point further in the future would not still be in the future.

Connectedness

These properties aren't enough to narrow down a unique interpretation of time, however. For one thing, we haven't made any claim that every time is comparable to every other, i.e. while we have said that not both $sR_M t$ and $tR_M s$ are true (a combination of antisymmetry and irreflexivity), it is possible that neither are true. This allows unconnected sections of the timeline, which have no relation to other sections. It also raises the possibility of branching time: there may be more than one possible future for a given point in time.

If we do make the requirement that exactly one of $sR_M t$ or $tR_M s$ are true, except in the case when $s = t$ (in which case both are false because of irreflexivity), then R_M fits most of our intuitions about time. It behaves like the "less than" relation on numbers, so we can consider the states of the model to each have a unique "clock".

Variations of temporal logic

Diodorean

If we instead allow R_M to be reflexive, we are dealing with the "Diodorean Modality" (Goldblatt pg. 44). This is used in an analysis of programs which we will see shortly. It forces us to consider the present to be part of the future, so what will always be true must currently be true.

Two-way

Often it is the case that we want to talk about times that are in the past as well as times that are in the future. We can do this by introducing another relation that is the converse (every pair is reversed) of the relation we have dealt with so far. We then also have two sets of modal connectives, instead of just \diamond and \square . These are $[F]$ and $\langle F \rangle$ for the future, and $[P]$ and $\langle P \rangle$ for the past.

Discrete vs continuous

In an analysis of computation we probably want to think of time as a series of instants, one following another. But in more general analyses of time, it may be desirable to assume that between any two distinct instants there is another.

There are several possible models of continuous time. Time could behave like the rational numbers, which can be indexed by the integers, or like the real numbers, which cannot. The distinction between real and rational time is only relevant with if we are using the two-way temporal logic described above; with only a future modality, they behave identically (Goldblatt ex. 6.4 and 8.8).

Dynamic logic

This section is largely drawn from Harel et al. Chapter 5.

A logic for change

Dynamic logic is about expressing the means of change between the states of a system, and how truths can be related between the states. Typically we think of programs changing the state of a computer. There are many (infinitely many) relations in dynamic logic, and each one represents the behavior of a particular program. These relations are completely arbitrary, in particular a program may relate one input state to several output states, which represents nondeterminism (i.e. multiple possible outcomes). We may use nondeterminism because there are truly unpredictable behaviors at work, or because it is preferable to not overspecify exactly what will happen.

A set of arbitrary relations is not a very effective representation of programming: while a program can be thought of as some mapping of inputs to outputs, it is usually constructed from a few primitive operations combined into a larger program. Programming languages provide constructs like loops and conditional statements to facilitate combining smaller parts in well-structured ways. Dynamic logic supports this view by providing ways to construct programs by combining other programs and by incorporating propositions.

Halting

If we expect an output from a program, we require that program to *halt*, i.e. to finish with a state we consider to be the output. Programs may not halt, which may be considered as reaching some error condition or simply running forever without producing a result. Dynamic logic inherently supports this behavior. Failing to halt on some input is indicated by a program relation mapping some input state to zero output states.

Syntax

At its heart, a dynamic logic is based on some set of primitive programs, these will just be the relations of the modal logic. For example we will call two primitive programs a and b . Programs can be combined in several ways, and each combination is itself a program.

- A *sequence*, written $a ; b$, runs program a and then program b .
- A *choice*, written $a \cup b$, nondeterministically runs either a or b .
- An *iteration*, written a^* , runs program a a nondeterministic number of times, 0 or more.

When we want to make propositions about what is true after running a program, we write the program as a modal connective, like the modal connectives \Box and \Diamond that we had in the basic modal logic language discussed

earlier. For instance we write $\langle a \rangle$ (like \diamond) if we are stating what is true after *some* run of program a , or $[a]$ (like \Box) if we are stating what is true after *all* runs of the program.

Thus we can make statements like $\varphi \rightarrow [a; b^*]\varphi$, meaning "if φ is true, then after running a and then running b any number of times, φ will be true". One interesting proposition is $[\alpha^*]\varphi \rightarrow [\alpha; \alpha^*]\varphi$, which states that if running α any number of times results in φ true, then running it at least once (once plus any number of times) results in φ true, this is true of any program α and any proposition φ .

As an example of choice, $\langle a \cup (b; b) \rangle \varphi$ means that if we consider running program a once or program b twice, at least one of these will result in φ .

We will usually use the $[a]$ form, as we are interested in what will necessarily be true after the program halts, but as noted we need to be careful that the program actually does halt.

Tests

There is one other way to make a program in dynamic logic: If φ is a proposition, $\varphi?$ is a program. This type of program is a *test*, and it can be understood as doing nothing when the proposition is true, thus the relation for $\varphi?$ relates input states where φ is true to identical outputs. However, if the proposition is false, the relation maps to no outputs, thus the program does not halt. This means that $[\varphi?]\varphi$ is always true; either

- φ was true and thus it is still true after $\varphi?$ (as it relates to the same output state), or
- φ was false so $\varphi?$ doesn't halt (and thus relates to no output state, so $[\varphi?]\psi$ is true for any ψ)

Consider a program that runs the primitive program a only if proposition p is true first, otherwise it does nothing. Perhaps it is not necessary or safe to run a if p is false. We can write the true part easily enough as $p? ; a$, which will halt with the result of running a if p was true (assuming a halts), but this will not halt if p is false. On the other hand $\neg p?$ will halt when p is false, without changing the state, but it will not halt when p is true. We can make this dilemma work with a choice: $(p? ; a) \cup \neg p?$. The \cup makes a program that tries one of the two branches, and as we have established only one of them will halt for a given truth value of p . As the choice is nondeterministic, the combined program will always halt, having run one of the two possibilities.

Formal semantics

Formally we provide the following rules for determining the truth of propositions in dynamic logic:

$M \models_s p$	if p is in $V_M(s)$
$M \models_s \neg \varphi$	if $M \not\models_s \varphi$
$M \models_s \varphi \rightarrow \psi$	if $M \models_s \varphi$ implies $M \models_s \psi$
$M \models_s \langle \alpha \rangle \varphi$	if there is some t such that $s(R\alpha)t$ and $M \models_t \varphi$
$M \models_s [\alpha] \varphi$	if for all t , if $s(R\alpha)t$, then $M \models_t \varphi$

(Note: \vee and \wedge can be defined in terms of \rightarrow and \neg).

$R\pi$ is the relation for a program π . $R\pi$ is defined for all non-primitive programs with the following rules, where α and β are programs:

$s(R\alpha; \beta)t$	if for some u , $s(R\alpha)u$ and $u(R\beta)t$
$s(R\alpha \cup \beta)t$	if $sR\alpha t$ or $sR\beta t$
$s(R\alpha^*)t$	if $s = t$ or $s(R\alpha; \alpha^*)t$
$s(R\varphi?)s$	if $M \models_s \varphi$

While programs

It is possible to express real programs in the language of dynamic logic, but often we want to use something closer to the programming languages in common use. The "while" language is one such way of expressing programs more intuitively. We will show how while language programs can be converted into dynamic logic.

primitives

We start with primitive programs that work like those in dynamic logic, they make some change to the program state. These aren't defined any more specifically by the language, the particular problem situation must specify the primitive programs that all larger programs are constructed from.

sequence

As with dynamic logic, we can write a program that runs program x and then program y as $x ; y$. This syntax is identical to the equivalent dynamic logic program.

conditional

The while language has a conditional statement of the form *if p then x else y* . If the proposition p is true then the program x is executed, if instead p is false then the program y is executed. The dynamic logic program $(p?; x) \cup (\neg p?; y)$ will accomplish the same thing, we saw a similar example earlier when introducing tests.

loop

The main loop construct in a while program is, naturally, *while*. We write *while p do x* to test proposition p , if it is false then the loop is finished, otherwise we run x and repeat from the test. This is expressed in dynamic logic as $(p?; x)^*; \neg p?$. The iterative part $((p?; x)^*)$ will test p (not progressing if p is false) and run x , repeatedly. We then need the extra "guard" test $\neg p?$ to prevent any early exit from the loop; since the iterative part is nondeterministic, it could run 0 times or any number of times until p becomes false, we only want it to run exactly as many times as it takes for p to become false.

Bibliography

- Patrick Blackburn, Maarten de Rijke, and Yde Venema. 2001. Modal Logic. Cambridge University Press.
- Robert Goldblatt. 1992. Logics of Time and Computation. Center for Study of Lang. and Info.
- David Harel, Jerzy Tiuryn, and Dexter Kozen. 2000. Dynamic Logic. MIT Press,.
- B. Mishra. 2013 (unpublished). Logic In Computer Science lecture notes.
- Amir Pnueli. 1981. "The temporal semantics of concurrent programs". Theoretical Computer Science, vol. 13, no. 1, pp. 45–60.